

If you set the cursor in XAML, you don't need to use the `Cursors` class directly. That's because the `TypeConverter` for the `Cursor` property is able to recognize the property names and retrieve the corresponding `Cursor` object from the `Cursors` class. That means you can write markup like this to show the help cursor (a combination of an arrow and a question mark) when the mouse is positioned over a button:

```
<Button Cursor="Help">Help</Button>
```

It's possible to have overlapping cursor settings. In this case, the most specific cursor wins. For example, you could set a different cursor on a button and on the window that contains the button. The button's cursor will be shown when you move the mouse over the button, and the window's cursor will be used for every other region in the window.

However, there's one exception. A parent can override the cursor settings of its children using the `ForceCursor` property. When this property is set to true, the child's `Cursor` property is ignored, and the parent's `Cursor` property applies everywhere inside.

If you want to apply a cursor setting to every element in every window of an application, the `FrameworkElement.Cursor` property won't help you. Instead, you need to use the static `Mouse.OverrideCursor` property, which overrides the `Cursor` property of every element:

```
Mouse.OverrideCursor = Cursors.Wait;
```

To remove this application-wide cursor override, set the `Mouse.OverrideCursor` property to null.

Lastly, WPF supports custom cursors without any fuss. You can use both ordinary `.cur` cursor files (which are essentially small bitmaps) and `.ani` animated cursor files. To use a custom cursor, you pass the file name of your cursor file or a stream with the cursor data to the constructor of the `Cursor` object:

```
Cursor customCursor = new Cursor(Path.Combine(applicationDir, "stopwatch.ani");
this.Cursor = customCursor;
```

The `Cursor` object doesn't directly support the URI resource syntax that allows other WPF elements (such as the `Image`) to use files that are stored in your compiled assembly. However, it's still quite easy to add a cursor file to your application as a resource and then retrieve it as a stream that you can use to construct a `Cursor` object. The trick is using the `Application.GetResourceStream()` method:

```
StreamResourceInfo sri = Application.GetResourceStream(
    new Uri("stopwatch.ani", UriKind.Relative));
Cursor customCursor = new Cursor(sri.Stream);
this.Cursor = customCursor;
```

This code assumes that you've added a file named `stopwatch.ani` to your project and set its `Build Action` to `Resource`. You'll learn more about the `GetResourceStream()` method in Chapter 7.

Content Controls

A *content control* is a still more specialized type of controls that is able to hold (and display) a piece of content. Technically, a content control is a control that can contain a single nested element. The one-child limit is what differentiates content controls from layout containers, which can hold as many nested elements as you want.

■ **Tip** Of course, you can still pack a lot of content in a single content control. The trick is to wrap everything in a single container, such as a StackPanel or a Grid. For example, the Window class is itself a content control. Obviously, windows often hold a great deal of content, but it's all wrapped in one top-level container (typically a Grid).

As you learned in Chapter 3, all WPF layout containers derive from the abstract Panel class, which gives the support for holding multiple elements. Similarly, all content controls derive from the abstract ContentControl class. Figure 6-1 shows the class hierarchy.




Figure 6-1. The hierarchy of content controls

As Figure 6-1 shows, several common controls are actually content controls, including the Label and the ToolTip. Additionally, all types of buttons are content controls, including the familiar Button, the RadioButton, and the CheckBox. There are also a few more specialized content controls, such as ScrollViewer (which allows you to create a scrollable panel) and UserControl class (which allows you to reuse a custom grouping of controls). The Window class, which is used to represent each window in your application, is itself a content control.

Finally, there is a subset of content controls that goes through one more level of inheritance by deriving from the HeaderedContentControl class. These controls have both a content region and a header region, which can be used to display some sort of title. They include the GroupBox, TabItem (a page in a TabControl), and Expander controls.

■ **Note** Figure 6-1 leaves out just a few elements. It doesn't show the Frame element, which is used for navigation (discussed in Chapter 24), and it omits a few elements that are used inside other controls (such as list box and status bar items).

The Content Property

Whereas the Panel class adds the Children collection to hold nested elements, the ContentControl class adds a Content property, which accepts a single object. The Content property supports any type of object, but it separates objects into two groups and gives each group different treatment:

- **Objects that don't derive from UIElement.** The content control calls ToString() to get the text for these controls and then displays that text.
- **Objects that derive from UIElement.** These objects (which include all the visual elements that are a part of WPF) are displayed inside the content control using the UIElement.OnRender() method.

■ **Note** Technically, the OnRender() method doesn't draw the object immediately. It simply generates a graphical representation, which WPF paints on the screen as needed.

To understand how this works, consider the humble button. So far, the examples that you've seen that include buttons have simply supplied a string:

```
<Button Margin="3">Text content</Button>
```

This string is set as the button content and displayed on the button surface. However, you can get more ambitious by placing other elements inside the button. For example, you can place an image inside a button using the Image class:

```
<Button Margin="3">
  <Image Source="happyface.jpg" Stretch="None" />
</Button>
```

Or you could combine text and images by wrapping them all in a layout container like the `StackPanel`:

```
<Button Margin="3">
  <StackPanel>
    <TextBlock Margin="3">Image and text button</TextBlock>
    <Image Source="happyface.jpg" Stretch="None" />
    <TextBlock Margin="3">Courtesy of the StackPanel</TextBlock>
  </StackPanel>
</Button>
```

■ **Note** It's acceptable to place text content inside a content control because the XAML parser converts that to a string object and uses that to set the `Content` property. However, you can't place string content directly in a layout container. Instead, you need to wrap it in a class that derives from `UIElement`, such as `TextBlock` or `Label`.

If you wanted to create a truly exotic button, you could even place other content controls such as text boxes and buttons inside the button (and still nest elements inside these). It's doubtful that such an interface would make much sense, but it's possible. Figure 6-2 shows some sample buttons.

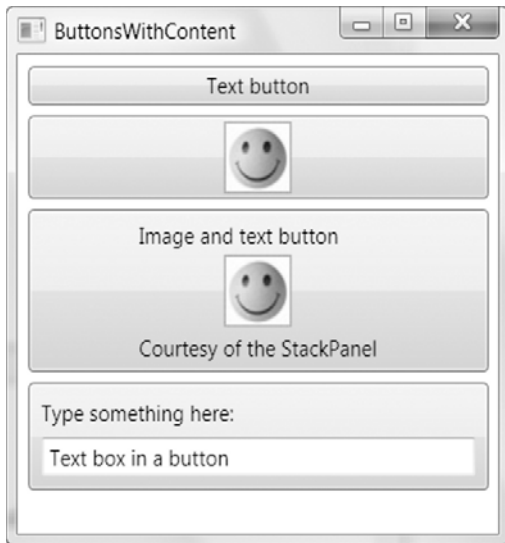


Figure 6-2. Buttons with different types of nested content

This is the same content model you saw with windows. Just like the `Button` class, the `Window` class allows a single nested element, which can be a piece of text, an arbitrary object, or an element.

■ **Note** One of the few elements that is *not* allowed inside a content control is the `Window`. When you create a `Window`, it checks to see if it's the top-level container. If it's placed inside another element, the `Window` throws an exception.

Aside from the `Content` property, the `ContentControl` class adds very little. It includes a `HasContent` property that returns true if there is content in the control, and a `ContentTemplate` that allows you to build a template telling the control how to display an otherwise unrecognized object. Using a `ContentTemplate`, you can display non-`UIElement`-derived objects more intelligently. Instead of just calling `ToString()` to get a string, you can take various property values and arrange them into more complex markup. You'll learn more about data templates in Chapter 20.

Aligning Content

In Chapter 3, you learned how to align different controls in a container using the `HorizontalAlignment` and `VerticalAlignment` properties, which are defined in the base `FrameworkElement` class. However, once a control contains content, you need to consider another level of organization. You need to decide how the content inside your content control is aligned with its borders. This is accomplished using the `HorizontalContentAlignment` and `VerticalContentAlignment` properties.

`HorizontalContentAlignment` and `VerticalContentAlignment` support the same values as `HorizontalAlignment` and `VerticalAlignment`. That means you can line up content on the inside of any edge (Top, Bottom, Left, or Right), you can center it (Center), or you can stretch it to fill the available space (Stretch). These settings are applied directly to the nested content element, but you can use multiple levels of nesting to create a sophisticated layout. For example, if you nest a `StackPanel` in a `Label` element, the `Label.HorizontalContentAlignment` property determines where the `StackPanel` is placed, but the alignment and sizing options of the `StackPanel` and its children will determine the rest of the layout.

In Chapter 3, you also learned about the `Margin` property, which allows you to add whitespace between adjacent elements. Content controls use a complementary property named `Padding`, which inserts space between the edges of the control and the edges of the content. To see the difference, compare the following two buttons:

```
<Button>Absolutely No Padding</Button>
<Button Padding="3">Well Padded</Button>
```

The button that has no padding (the default) has its text crowded against the button edge. The button that has a padding of 3 units on each side gets a more respectable amount of breathing space. Figure 6-3 highlights the difference.



Figure 6-3. *Padding the content of the button*

■ **Note** The `HorizontalAlignment`, `VerticalContentAlignment`, and `Padding` properties are defined as part of the `Control` class, not the more specific `ContentControl` class. That's because there may be controls that aren't content controls but still have some sort of content. One example is the `TextBox`—its contained text (stored in the `Text` property) is adjusted using the alignment and padding settings you've applied.

The WPF Content Philosophy

At this point, you might be wondering if the WPF content model is really worth all the trouble. After all, you might choose to place an image inside a button, but you're unlikely to embed other controls and entire layout panels. However, there are a few important reasons driving the shift in perspective.

Consider the example shown in Figure 6-2, which includes a simple image button that places an `Image` element inside the `Button` control. This approach is less than ideal, because bitmaps are not resolution-independent. On a high-dpi display, the bitmap may appear blurry because WPF must add more pixels by interpolation to make sure the image stays the correct size. More sophisticated WPF interfaces avoid bitmaps and use a combination of vector shapes to create custom-drawn buttons and other graphical frills (as you'll see in Chapter 12).

This approach integrates nicely with the content control model. Because the `Button` class is a content control, you are not limited to filling it with a fixed bitmap; instead, you can include other content. For example, you can use the classes in the `System.Windows.Shapes` namespace to draw a vector image inside a button. Here's an example that creates a button with two diamond shapes (as shown in Figure 6-4):

```
<Button Margin="3">
  <Grid>
    <Polygon Points="100,25 125,0 200,25 125,50"
      Fill="LightSteelBlue" />
    <Polygon Points="100,25 75,0 0,25 75,50"
      Fill="White"/>
  </Grid>
</Button>
```




Figure 6-4. A button with shape content

Clearly, in this case, the nested content model is simpler than adding extra properties to the Button class to support the different types of content. Not only is the nested content model more flexible, but it also allows the Button class to expose a simpler interface. And because all content controls support content nesting in the same way, there's no need to add different content properties to multiple classes. (Windows Forms ran into this issue in .NET 2.0, while enhancing the Button and Label class to better support images and mixed image and text content.)

In essence, the nested content model is a trade-off. It simplifies the class model for elements because there's no need to use additional layers of inheritance to add properties for different types of content. However, you need to use a slightly more complex *object* model—elements that can be built from other nested elements.

■ **Note** You can't always get the effect you want by changing the content of a control. For example, even though you can place any content in a button, a few details never change, such as the button's shaded background, its rounded border, and the mouse-over effect that makes it glow when you move the mouse pointer over it. However, another way to change these built-in details is to apply a new control template. Chapter 17 shows how you can change all aspects of a control's look and feel using a control template.

Labels

The simplest of all content controls is the Label control. Like any other content control, it accepts any single piece of content you want to place inside. But what distinguishes the Label control is its support for *mnemonics*, which are essentially shortcut keys that set the focus to a linked control.

To support this functionality, the Label control adds a single property, named *Target*. To set the Target property, you need to use a binding expression that points to another control. Here's the syntax you must use:

```
<Label Target="{Binding ElementName=txtA}">Choose _A</Label>
<TextBox Name="txtA"></TextBox>
<Label Target="{Binding ElementName=txtB}">Choose _B</Label>
<TextBox Name="txtB"></TextBox>
```

The underscore in the label text indicates the shortcut key. (If you really *do* want an underscore to appear in your label, you must add two underscores instead.) All mnemonics work with Alt and the shortcut key you've identified. For example, if the user presses Alt+A in this example, the first label transfers focus to the linked control, which is txtA. Similarly, Alt+B takes the user to txtB.

■ **Note** If you've programmed with Windows Forms, you're probably used to using the ampersand (&) character to identify a shortcut key. XAML uses the underscore instead because the ampersand character can't be entered directly in XML; instead, you need to use the clunkier character entity &#amp; in its place.

Usually, the shortcut letters are hidden until the user presses Alt, at which point they appear as underlined letters (Figure 6-5). However, this behavior depends on system settings.

■ **Tip** If all you need to do is display content without support for mnemonics, you may prefer to use the more lightweight TextBlock element. Unlike the Label, the TextBlock also supports wrapping through its TextWrapping property.




Figure 6-5. Shortcuts in a label

Buttons

WPF recognizes three types of button controls: the familiar `Button`, the `CheckBox`, and the `RadioButton`. All of these controls are content controls that derive from `ButtonBase`.

The `ButtonBase` class includes only a few members. It defines the `Click` event and adds support for commands, which allow you to wire buttons to higher-level application tasks (a feat you'll consider in Chapter 9). Finally, the `ButtonBase` class adds a `ClickMode` property, which determines when a button fires its `Click` event in response to mouse actions. The default value is `ClickMode.Release`, which means the `Click` event fires when the mouse is clicked and released. However, you can also choose to fire the `Click` event mouse when the mouse button is first pressed (`ClickMode.Press`) or, oddly enough, whenever the mouse moves over the button and pauses there (`ClickMode.Hover`).

■ **Note** All button controls support access keys, which work similarly to mnemonics in the `Label` control. You add the underscore character to identify the access key. If the user presses `Alt` and the access key, a button click is triggered.

The Button

The `Button` class represents the ever-present Windows push button. It adds just two writeable properties, `IsCancel` and `IsDefault`:

- **When `IsCancel` is true**, this button is designated as the cancel button for a window. If you press the `Escape` key while positioned anywhere on the current window, this button is triggered.
- **When `IsDefault` is true**, this button is designated as the default button (also known as the accept button). Its behavior depends on your current location in the window. If you're positioned on a non-`Button` control (such as a `TextBox`, `RadioButton`, `CheckBox`, and so on), the default button is given a blue shading, almost as though it has focus. If you press `Enter`, this button is triggered. However, if you're positioned on another `Button` control, the current button gets the blue shading, and pressing `Enter` triggers that button, not the default button.

Many users rely on these shortcuts (particularly the `Escape` key to close an unwanted dialog box), so it makes sense to take the time to define these details in every window you create. It's still up to you to write the event handling code for the cancel and default buttons, because WPF won't supply this behavior.

In some cases, it may make sense for the same button to be the cancel button *and* the default button for a window. One example is the `OK` button in an `About` box. However, there should be only a single cancel button and a single default button in a window. If you designate more than one cancel button, pressing `Escape` will simply move the focus to the next default button, but it won't trigger that button. If you have more than one default button, pressing `Enter` has a somewhat more confusing behavior. If you're on a non-`Button` control, pressing `Enter` moves you to the next default button. If you're on a `Button` control, pressing `Enter` triggers it.

IsDefault and IsDefaulted

The `Button` class also includes the horribly confusing `IsDefaulted` property, which is read-only. `IsDefaulted` returns true for a default button if another control has focus and that control doesn't accept the Enter key. In this situation, pressing the Enter key will trigger the button.

For example, a `TextBox` does not accept the Enter key, unless you've set `TextBox.AcceptsReturn` to true. When a `TextBox` with an `AcceptsReturn` value of true has focus, `IsDefaulted` is false for the default button. When a `TextBox` with an `AcceptsReturn` value of false has focus, the default button has `IsDefaulted` set to true. If this isn't confusing enough, the `IsDefaulted` property returns false when the button itself has focus, even though hitting Enter at this point will trigger the button.

Although it's unlikely that you'll want to use the `IsDefaulted` property, this property does allow you to write certain types of style triggers, as you'll see in Chapter 11. If that doesn't interest you, just add it to your list of obscure WPF trivia, which you can use to puzzle your colleagues.

The ToggleButton and RepeatButton

Along with `Button`, three more classes derive from `ButtonBase`. These include the following:

- `GridViewColumnHeader`, which represents the clickable header of a column when you use a grid-based `ListView`. The `ListView` is described in Chapter 22.
- `RepeatButton`, which fires Click events continuously, as long as the button is held down. Ordinary buttons fire one Click event per user click.
- `ToggleButton`, which represents a button that has two states (pushed or unpushed). When you click a `ToggleButton`, it stays in its pushed state until you click it again to release it. This is sometimes described as *sticky click* behavior.

Both `RepeatButton` and `ToggleButton` are defined in the `System.Windows.Controls.Primitives` namespace, which indicates they aren't often used on their own. Instead, they're used to build more complex controls by composition, or extended with features through inheritance. For example, the `RepeatButton` is used to build the higher-level `ScrollBar` control (which, ultimately, is a part of the even higher-level `ScrollViewer`). The `RepeatButton` gives the arrow buttons at the ends of the scroll bar their trademark behavior—scrolling continues as long as you hold it down. Similarly, the `ToggleButton` is used to derive the more useful `CheckBox` and `RadioButton` classes described next.

However, neither the `RepeatButton` nor the `ToggleButton` is an abstract class, so you can use both of them directly in your user interfaces. The `ToggleButton` is genuinely useful inside a `ToolBar`, which you'll use in Chapter 25.

The CheckBox

Both the `CheckBox` and the `RadioButton` are buttons of a different sort. They derive from `ToggleButton`, which means they can be switched on or off by the user, hence their “toggle” behavior. In the case of the `CheckBox`, switching the control on means placing a check mark in it.

The `CheckBox` class doesn't add any members, so the basic `CheckBox` interface is defined in the `ToggleButton` class. Most important, `ToggleButton` adds an `IsChecked` property. `IsChecked` is a nullable `Boolean`, which means it can be set to true, false, or null. Obviously, true represents a checked box, while false represents an empty one. The null value is a little trickier—it represents an indeterminate state, which is displayed as a shaded box. The indeterminate state is commonly used to represent values that haven't been set or areas where some discrepancy exists. For example, if you have a check box that allows you to apply bold formatting in a text application and the current selection includes both bold and regular text, you might set the check box to null to show an indeterminate state.

To assign a null value in WPF markup, you need to use the null markup extension, as shown here:

```
<CheckBox IsChecked="{x:Null}">A check box in indeterminate state</CheckBox>
```

Along with the `IsChecked` property, the `ToggleButton` class adds a property named `IsThreeState`, which determines whether the user is able to place the check box into an indeterminate state. If `IsThreeState` is false (the default), clicking the check box alternates its state between checked and unchecked, and the only way to place it in an indeterminate state is through code. If `IsThreeState` is true, clicking the check box cycles through all three possible states.

The `ToggleButton` class also defines three events that fire when the check box enters specific states: `Checked`, `Unchecked`, and `Indeterminate`. In most cases, it's easier to consolidate this logic into one event handler by handling the `Click` event that's inherited from `ButtonBase`. The `Click` event fires whenever the button changes state.

The RadioButton

The `RadioButton` also derives from `ToggleButton` and uses the same `IsChecked` property and the same `Checked`, `Unchecked`, and `Indeterminate` events. Along with these, the `RadioButton` adds a single property named `GroupName`, which allows you to control how radio buttons are placed into groups.

Ordinarily, radio buttons are grouped by their container. That means if you place three `RadioButton` controls in a single `StackPanel`, they form a group from which you can select just one of the three. On the other hand, if you place a combination of radio buttons in two separate `StackPanel` controls, you have two independent groups on your hands.

The `GroupName` property allows you to override this behavior. You can use it to create more than one group in the same container or to create a single group that spans multiple containers. Either way, the trick is simple—just give all the radio buttons that belong together the same group name.

Consider this example:

```
<StackPanel>
  <GroupBox Margin="5">
    <StackPanel>
      <RadioButton>Group 1</RadioButton>
      <RadioButton>Group 1</RadioButton>
      <RadioButton>Group 1</RadioButton>
      <RadioButton Margin="0,10,0,0" GroupName="Group2">Group 2</RadioButton>
    </StackPanel>
  </GroupBox>
  <GroupBox Margin="5">
    <StackPanel>
      <RadioButton>Group 3</RadioButton>
      <RadioButton>Group 3</RadioButton>
      <RadioButton>Group 3</RadioButton>
      <RadioButton Margin="0,10,0,0" GroupName="Group2">Group 2</RadioButton>
    </StackPanel>
  </GroupBox>
</StackPanel>
```

```

    </StackPanel>
  </GroupBox>
</StackPanel>

```

Here, there are two containers holding radio buttons, but three groups. The final radio button at the bottom of each group box is part of a third group. In this example, it makes for a confusing design, but there may be some scenarios where you want to separate a specific radio button from the pack in a subtle way without causing it to lose its group membership.

■ **Tip** You don't need to use the `GroupBox` container to wrap your radio buttons, but it's a common convention. The `GroupBox` shows a border and gives you a caption that you can apply to your group of buttons.

Tooltips

WPF has a flexible model for *tooltips* (those infamous yellow boxes that pop up when you hover over something interesting). Because tooltips in WPF are content controls, you can place virtually anything inside a tooltip. You can also tweak various timing settings to control how quickly tooltips appear and disappear.

The easiest way to show a tooltip doesn't involve using the `ToolTip` class directly. Instead, you simply set the `ToolTip` property of your element. The `ToolTip` property is defined in the `FrameworkElement` class, so it's available on anything you'll place in a WPF window.

For example, here's a button that has a basic tooltip:

```
<Button ToolTip="This is my tooltip">I have a tooltip</Button>
```

When you hover over this button, the text “This is my tooltip” appears in the familiar yellow box.

If you want to supply more ambitious tooltip content, such as a combination of nested elements, you need to break the `ToolTip` property out into a separate element. Here's an example that sets the `ToolTip` property of a button using more complex nested content:

```

<Button>
  <Button.ToolTip>
    <StackPanel>
      <TextBlock Margin="3" >Image and text</TextBlock>
      <Image Source="happyface.jpg" Stretch="None" />
      <TextBlock Margin="3" >Image and text</TextBlock>
    </StackPanel>
  </Button.ToolTip>
  <Button.Content>I have a fancy tooltip</Button.Content>
</Button>

```

As in the previous example, WPF implicitly creates a `ToolTip` object. The difference is that, in this case, the `ToolTip` object contains a `StackPanel` rather than a simple string. Figure 6-6 shows the result.




Figure 6-6. *A fancy tooltip*

If more than one tooltip overlaps, the most specific tooltip wins. For example, if you add a tooltip to the StackPanel container in the previous example, this tooltip appears when you hover over an empty part of the panel or a control that doesn't have its own tooltip.

■ **Note** Don't put user-interactive controls in a tooltip because the ToolTip window can't accept focus. For example, if you place a button in a ToolTip control, the button will appear, but it isn't clickable. (If you attempt to click it, your mouse click will just pass through to the window underneath.) If you want a tooltip-like window that can hold other controls, consider using the Popup control instead, which is discussed shortly, in the section named "The Popup."

Setting ToolTip Properties

The previous example shows how you can customize the content of a tooltip, but what if you want to configure other ToolTip-related settings? You actually have two options. The first technique you can use is to explicitly define the ToolTip object. That gives you the chance to directly set a variety of ToolTip properties.

The ToolTip is a content control, so you can adjust standard properties such as the Background (so it isn't a yellow box), Padding, and Font. You can also modify the members that are defined in the ToolTip class (listed in Table 6-2). Most of these properties are designed to help you place the tooltip exactly where you want it.

Table 6-2. *ToolTip Properties*

Name	Description
HasDropShadow	Determines whether the tooltip has a diffuse black drop shadow that makes it stand out from the window underneath.
Placement	Determines how the tooltip is positioned, using one of the values from the <code>PlacementMode</code> enumeration. The default value is <code>Mouse</code> , which means that the top-left corner of the tooltip is placed relative to the current mouse position. (The actual position of the tooltip may be offset from this starting point based on the <code>HorizontalOffset</code> and <code>VerticalOffset</code> properties.) Other possibilities allow you to place the tooltip using absolute screen coordinates or place it relative to some element (which you indicate using the <code>PlacementTarget</code> property).
HorizontalOffset and VerticalOffset	Allows you to nudge the tooltip into the exact position you want. You can use positive or negative values.
PlacementTarget	Allows you to place a tooltip relative to another element. In order to use this property, the <code>Placement</code> property must be set to <code>Left</code> , <code>Right</code> , <code>Top</code> , <code>Bottom</code> , or <code>Center</code> . (This is the edge of the element to which the tooltip is aligned.)
PlacementRectangle	Allows you to offset the position of the tooltip. This works in much the same way as the <code>HorizontalOffset</code> and <code>VerticalOffset</code> properties. This property doesn't have an effect if <code>Placement</code> property is set to <code>Mouse</code> .
CustomPopupPlacement Callback	Allows you to position a tooltip dynamically using code. If the <code>Placement</code> property is set to <code>Custom</code> , this property identifies the method that will be called by the <code>ToolTip</code> to get the position where the <code>ToolTip</code> should be placed. Your callback method receives three pieces of information: <code>popupSize</code> (the size of the <code>ToolTip</code>), <code>targetSize</code> (the size of the <code>PlacementTarget</code> , if it's used), and <code>offset</code> (a point that's created based on <code>HorizontalOffset</code> and <code>VerticalOffset</code> properties). The method returns a <code>CustomPopupPlacement</code> object that tells WPF where to place the tooltip.
StaysOpen	Has no effect in practice. The intended purpose of this property is to allow you to create a tooltip that remains open until the user clicks somewhere else. However, the <code>ToolTipService.ShowDuration</code> property overrides the <code>StaysOpen</code> property. As a result, tooltips always disappear after a configurable amount of time (usually about 5 seconds) or when the user moves the mouse away. If you want to create a tooltip-like window that stays open indefinitely, the easiest approach is to use the <code>Popup</code> control.
IsEnabled and IsOpen	Allow you to control the tooltip in code. <code>IsEnabled</code> allows you to temporarily disable a <code>ToolTip</code> . <code>IsOpen</code> allows you to programmatically show or hide a tooltip (or just check whether the tooltip is open).

Using the `ToolTip` properties, the following markup creates a tooltip that has no drop shadow but uses a transparent red background, which lets the underlying window (and controls) show through:

```
<Button>
  <Button.ToolTip>
    <ToolTip Background="#60AA4030" Foreground="White"
      HasDropShadow="False" >
      <StackPanel>
        <TextBlock Margin="3" >Image and text</TextBlock>
        <Image Source="happyface.jpg" Stretch="None" />
        <TextBlock Margin="3" >Image and text</TextBlock>
      </StackPanel>
    </ToolTip>
  </Button.ToolTip>
  <Button.Content>I have a fancy tooltip</Button.Content>
</Button>
```

In most cases, you'll be happy enough to use the standard tooltip placement, which puts it at the current mouse position. However, the various `ToolTip` properties give you many more options. Here are some strategies you can use to place a tooltip:

- **Based on the current position of the mouse.** This is the standard behavior, which relies on `Placement` being set to `Mouse`. The top-left corner of the tooltip box is lined up with the bottom-left corner of the invisible bounding box around the mouse pointer.
- **Based on the position of the moused-over element.** Set the `Placement` property to `Left`, `Right`, `Top`, `Bottom`, or `Center`, depending on the edge of the element you want to use. The top-left corner of the tooltip box will be lined up with that edge.
- **Based on the position of another element (or the window).** Set the `Placement` property in the same way you would if you were lining up the tooltip with the current element. (Use the value `Left`, `Right`, `Top`, `Bottom`, or `Center`.) Then choose the element by setting the `PlacementTarget` property. Remember to use the `{Binding ElementName=Name}` syntax to identify the element you want to use.
- **With an offset.** Use any of the strategies described previously, but set the `HorizontalOffset` and `VerticalOffset` properties to add a little extra space.
- **Using absolute coordinates.** Set `Placement` to `Absolute` and use the `HorizontalOffset` and `VerticalOffset` properties (or the `PlacementRectangle`) to set some space between the tooltip and the top-left corner of the window.
- **Using a calculation at runtime.** Set `Placement` to `Custom`. Set the `CustomPopupPlacementCallback` property to point to a method that you've created.

Figure 6-7 shows how different placement properties stack up. Note that when lining up a tooltip against an element along the tooltip's bottom or right edge, you'll end up with a bit of extra space. That's because of the way that the `ToolTip` measures its content.

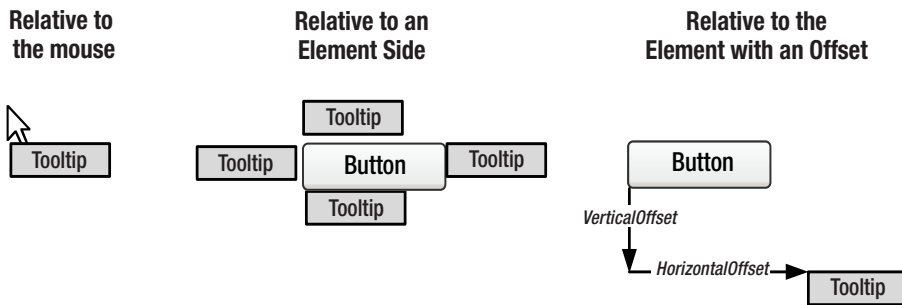


Figure 6-7. *Placing a tooltip explicitly*

Setting ToolTipService Properties

There are some tooltip properties that can't be configured using the properties of the `ToolTip` class. In this case, you need to use a different class, which is named `ToolTipService`. `ToolTipService` allows you to configure the time delays associated with the display of a tooltip. All the properties of the `ToolTipService` class are attached properties, so you can set them directly in your control tag, as shown here:

```
<Button ToolTipService.InitialShowDelay="1">
...
</Button>
```

The `ToolTipService` class defines many of the same properties as `ToolTip`. This allows you to use a simpler syntax when you're dealing with text-only tooltips. Rather than adding a nested `ToolTip` element, you can set everything you need using attributes:

```
<Button ToolTip="This tooltip is aligned with the bottom edge"
  ToolTipService.Placement="Bottom">I have a tooltip</Button>
```

Table 6-3 lists the properties of the `ToolTipService` class. The `ToolTipService` class also provides two routed events: `ToolTipOpening` and `ToolTipClosing`. You can react to these events to fill a tooltip with just-in-time content or to override the way tooltips work. For example, if you set the `handled` flag in both events, tooltips will no longer be shown or hidden automatically. Instead, you'll need to show and hide them manually by setting the `IsOpen` property.

■ **Tip** It makes little sense to duplicate the same tooltip settings for several controls. If you plan to adjust the way tooltips are handled in your entire application, use styles so that your settings are applied automatically, as described in Chapter 11. Unfortunately, the `ToolTipService` property values are not inherited, which means if you set them at the window or container level, they don't flow through to the nested elements.

Table 6-3. *ToolTipService Properties*

Name	Description
InitialShowDelay	Sets the delay (in milliseconds) before this tooltip is shown when the mouse hovers over the element.
ShowDuration	Sets the amount of time (in milliseconds) that this tooltip is shown before it disappears, if the user does not move the mouse.
BetweenShowDelay	Sets a time window (in milliseconds) during which the user can move between tooltips without experiencing the InitialShowDelay. For example, if BetweenShowDelay is 5000, the user has 5 seconds to move to another control that has a tooltip. If the user moves to another control within that time period, the new tooltip is shown immediately. If the user takes longer, the BetweenShowDelay window expires, and the InitialShowDelay kicks into action. In this case, the second tooltip isn't shown until after the InitialShowDelay period.
ToolTip	Sets the content for the tooltip. Setting ToolTipService.ToolTip is equivalent to setting the FrameworkElement.ToolTip property of an element.
HasDropShadow	Determines whether the tooltip has a diffuse black drop shadow that makes it stand out from the window underneath.
ShowOnDisabled	Determines the tooltip behavior when the associated element is disabled. If true, the tooltip will appear for disabled elements (elements that have their IsEnabled property set to false). The default is false, in which case the tooltip appears only if the associated element is enabled.
Placement, PlacementTarget, PlacementRectangle, and VerticalOffset	Allows you to control the placement of the tooltip. These properties work in the same way as the matching properties of the ToolTipHorizontalOffset class.

The Popup

The Popup control has a great deal in common with the ToolTip, although neither one derives from the other.

Like the ToolTip, the Popup can hold a single piece of content, which can include any WPF element. (This content is stored in the Popup.Child property, unlike the ToolTip content, which is stored in the ToolTip.Content property.) Also, like the ToolTip, the content in the Popup can extend beyond the bounds of the window. Lastly, the Popup can be placed using the same placement properties and shown or hidden using the same IsOpen property.

The differences between the Popup and ToolTip are more important. They include the following:

- The Popup is never shown automatically. You must set the `IsOpen` property for it to appear.
- By default, the `Popup.StaysOpen` property is set to `true`, and the Popup does not disappear until you explicitly set its `IsOpen` property to `false`. If you set `StaysOpen` to `false`, the Popup disappears when the user clicks somewhere else.

■ **Note** A popup that stays open can be a bit jarring because it behaves like a separate stand-alone window. If you move the window underneath, the popup remains fixed in its original position. You won't witness this behavior with the ToolTip or with a Popup that sets `StaysOpen` to `false`, because as soon as you click to move the window, the tooltip or popup window disappears.

- The Popup provides a `PopupAnimation` property that lets you control how it comes into view when you set `IsOpen` to `true`. Your options include `None` (the default), `Fade` (the opacity of the popup gradually increases), `Scroll` (the popup slides in from the upper-left corner of the window, space permitting), and `Slide` (the popup slides down into place, space permitting). In order for any of these animations to work, you must also set the `AllowsTransparency` property to `true`.
- The Popup can accept focus. Thus, you can place user-interactive controls in it, such as a Button. This functionality is one of the key reasons to use the Popup instead of the ToolTip.
- The Popup control is defined in the `System.Windows.Controls.Primitives` namespace because it is most commonly used as a building block for more complex controls. You'll find that the Popup is not quite as polished as other controls. Notably, you must set the `Background` property if you want to see your content, because it won't be inherited from your window and you need to add the border yourself (the `Border` element works perfectly well for this purpose).

Because the Popup must be shown manually, you may choose to create it entirely in code. However, you can define it just as easily in XAML markup—just make sure to include the `Name` property so you can manipulate it in code.

Figure 6-8 shows an example. Here, when the user moves the mouse over an underlined word, a popup appears with more information and a link that opens an external web browser window.




Figure 6-8. A popup with a hyperlink

To create this window, you need to include a `TextBlock` with the initial text and a `Popup` with the additional content that you'll show when the user moves the mouse into the correct place. Technically, it doesn't matter where you define the `Popup` tag, because it's not associated with any particular control. Instead, it's up to you to set the placement properties to position the `Popup` in the correct spot. In this example, the `Popup` appears at the current mouse position, which is the simplest option.

```
<TextBlock TextWrapping="Wrap">You can use a Popup to provide a link for a
specific <Run TextDecorations="Underline" MouseEnter="run_MouseEnter">term</Run>
of interest.</TextBlock>

<Popup Name="popLink" StaysOpen="False" Placement="Mouse" MaxWidth="200"
  PopupAnimation="Slide" AllowsTransparency="True">
  <Border BorderBrush="Beige" BorderThickness="2" Background="White">
    <TextBlock Margin="10" TextWrapping="Wrap">
      For more information, see
      <Hyperlink NavigateUri="http://en.wikipedia.org/wiki/Term"
        Click="lnk_Click">Wikipedia</Hyperlink>
    </TextBlock>
  </Border>
</Popup>
```

This example presents two elements that you might not have seen before. The `Run` element allows you to apply formatting to a specific part of a `TextBlock`—it's a piece of flow content that you'll learn about in Chapter 28 when you consider documents. The `Hyperlink` allows you to provide a clickable piece of text. You'll take a closer look at it in Chapter 24, when you consider page-based applications.

The only remaining details are the relatively trivial code that shows the Popup when the mouse moves over the correct word and the code that launches the web browser when the link is clicked:

```
private void run_MouseEnter(object sender, MouseEventArgs e)
{
    popLink.IsOpen = true;
}

private void lnk_Click(object sender, RoutedEventArgs e)
{
    Process.Start(((Hyperlink)sender).NavigateUri.ToString());
}
```

■ **Note** You can show and hide a Popup using a *trigger*—an action that takes place automatically when a specific property hits a specific value. You simply need to create a trigger that reacts when the `Popup.IsMouseOver` is true and sets the `Popup.IsOpen` property to true. Chapter 11 has the details.

Specialized Containers

Content controls aren't just for basics like labels, buttons, and tooltips. They also include specialized containers that allow you to shape large portions of your user interface. In the following sections, you'll meet some of these more ambitious content controls: the `ScrollViewer`, `GroupBox`, `TabItem`, and `Expander`. All of these controls are designed to help you shape large portions of your user interface. However, because these controls can hold only a single element, you'll usually use them in conjunction with a layout container.

The ScrollViewer

Scrolling is a key feature if you want to fit large amounts of content in a limited amount of space. In order to get scrolling support in WPF, you need to wrap the content you want to scroll inside a `ScrollViewer`.

Although the `ScrollViewer` can hold anything, you'll typically use it to wrap a layout container. For example, in Chapter 3, you saw an example that used a `Grid` element to create a three-column display of texts, text boxes, and buttons. To make this `Grid` scrollable, you simply need to wrap the `Grid` in a `ScrollViewer`, as shown in this slightly shortened markup:

```
<ScrollViewer>
  <Grid Margin="3,3,10,3">
    <Grid.RowDefinitions>
      ...
    </Grid.RowDefinitions>
    <Grid.ColumnDefinitions>
      ...
    </Grid.ColumnDefinitions>
```