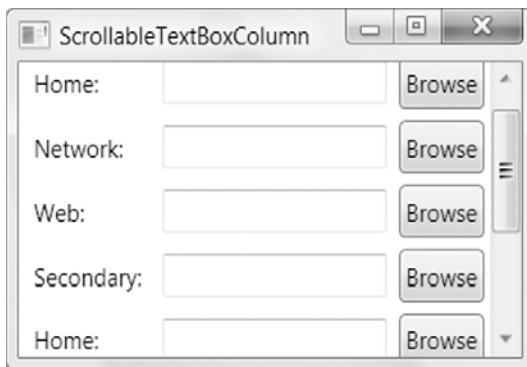


```

<Label Grid.Row="0" Grid.Column="0" Margin="3"
    VerticalAlignment="Center">Home:</Label>
<TextBox Grid.Row="0" Grid.Column="1" Margin="3"
    Height="Auto" VerticalAlignment="Center"></TextBox>
<Button Grid.Row="0" Grid.Column="2" Margin="3" Padding="2">
    Browse</Button>
...
</Grid>
</ScrollView>

```

The result is shown in Figure 6-9.



**Figure 6-9.** A scrollable window

If you resize the window in this example so that it's large enough to fit all its content, the scroll bar becomes disabled. However, the scroll bar will still be visible. You can control this behavior by setting the `VerticalScrollBarVisibility` property, which takes a value from the `ScrollBarVisibility` enumeration. The default value of `Visible` makes sure the vertical scroll bar is always present. Use `Auto` if you want the scroll bar to appear when it's needed and disappear when it's not. Or use `Disabled` if you don't want the scroll bar to appear at all.

---

■ **Note** You can also use `Hidden`, which is similar to `Disabled` but subtly different. First, content with a hidden scroll bar is still scrollable. (For example, you can scroll through the content using the arrow keys.) Second, the content in a `ScrollView` is laid out differently. When you use `Disabled`, you tell the content in the `ScrollView` that it has only as much space as the `ScrollView` itself. On the other hand, if you use `Hidden`, you tell the content that it has an infinite amount of space. That means it can overflow and stretch off into the scrollable region. Ordinarily, you'll use `Hidden` only if you plan to allow scrolling by another mechanism (such as the custom scrolling buttons described next). You'll use `Disabled` only if you want to temporarily prevent the `ScrollView` from doing anything at all.

---


The ScrollViewer also supports horizontal scrolling. However, the `HorizontalScrollBarVisibility` property is `Hidden` by default. To use horizontal scrolling, you need to change this value to `Visible` or `Auto`.

## Programmatic Scrolling

To scroll through the window shown in Figure 6-9, you can click the scroll bar with the mouse, you can move over the grid and use a mouse scroll wheel, you can tab through the controls, or you can click somewhere on the blank surface of the grid and use the up and down arrow keys. If this still doesn't give you the flexibility you crave, you can use the methods of the `ScrollViewer` class to scroll your content programmatically:

- The most obvious are `LineUp()` and `LineDown()`, which are equivalent to clicking the arrow buttons on the vertical scroll bar to move up or down once.
- You can also use `PageUp()` and `PageDown()`, which scroll an entire screenful up or down and are equivalent to clicking the surface of the scroll bar, above or below the scroll bar thumb.
- Similar methods allow horizontal scrolling, including `LineLeft()`, `LineRight()`, `PageLeft()`, and `PageRight()`.
- Finally, you can use the `ScrollToXxx()` methods to go somewhere specific. For vertical scrolling, they include `ScrollToEnd()` and `ScrollToHome()`, which take you to the top or bottom of the scrollable content, and `ScrollToVerticalOffset()`, which takes you to a specific position. There are horizontal versions of the same methods, including `ScrollToLeftEnd()`, `ScrollToRightEnd()`, and `ScrollToHorizontalOffset()`.

Figure 6-10 shows an example where several custom buttons allow you to move through the `ScrollViewer`. Each button triggers a simple event handler that uses one of the methods in the previous list.



**Figure 6-10.** Programmatic scrolling

## Custom Scrolling

The built-in scrolling in the `ScrollView` is quite useful. It allows you to scroll slowly through any content, from a complex vector drawing to a grid of elements. However, one of the most intriguing features of the `ScrollView` is its ability to let its content participate in the scrolling process. Here's how it works:

- You place a scrollable element inside the `ScrollView`. This is any element that implements `IScrollInfo`.
- You tell the `ScrollView` that the content knows how to scroll itself by setting the `ScrollView.CanContentScroll` property to `true`.
- When you interact with the `ScrollView` (by using the scroll bar, the mouse wheel, the scrolling methods, and so on), the `ScrollView` calls the appropriate methods on your element using the `IScrollInfo` interface. The element then performs its own custom scrolling.

---

■ **Note** The `IScrollInfo` interface defines a set of methods that react to different scrolling actions. For example, it includes many of the scrolling methods exposed by the `ScrollView`, such as `LineUp()`, `LineDown()`, `PageUp()`, and `PageDown()`. It also defines methods that handle the mouse wheel.

---

Very few elements implement `IScrollInfo`. One element that does is the `StackPanel` container. Its `IScrollInfo` implementation uses *logical scrolling*, which is scrolling that moves from element to element, rather than from line to line.

If you place a `StackPanel` in a `ScrollView` and you don't set the `CanContentScroll` property, you get the ordinary behavior. Scrolling up and down moves you a few pixels at a time. However, if you set `CanContentScroll` to `true`, each time you click down, you scroll to the beginning of the next element:

```
<ScrollView CanContentScroll="True">
  <StackPanel>
    <Button Height="100">1</Button>
    <Button Height="100">2</Button>
    <Button Height="100">3</Button>
    <Button Height="100">4</Button>
  </StackPanel>
</ScrollView>
```

You may or may not find that the `StackPanel`'s logical scrolling system is useful in your application. However, it's indispensable if you want to create a custom panel with specialized scrolling behavior.

## Headered Content Controls


One of the classes that derive from `ContentControl` is `HeaderedContentControl`. Its role is simple—it represents a container that has both single-element content (as stored in the `Content` property) and a single-element header (as stored in the `Header` property). The addition of the header is what distinguishes the `HeaderedContentControl` from the content controls you’ve seen so far.

Three classes derive from `HeaderedContentControl`: `GroupBox`, `TabItem`, and `Expander`. You’ll explore them in the following sections.

### The GroupBox

The `GroupBox` is the simplest of the three controls that derives from `HeaderedContentControl`. It’s displayed as a box with rounded corners and a title. Here’s an example (shown in Figure 6-11):

```
<GroupBox Header="A GroupBox Test" Padding="5"
  Margin="5" VerticalAlignment="Top">
  <StackPanel>
    <RadioButton Margin="3">One</RadioButton>
    <RadioButton Margin="3">Two</RadioButton>
    <RadioButton Margin="3">Three</RadioButton>
    <Button Margin="3">Save</Button>
  </StackPanel>
</GroupBox>
```



**Figure 6-11.** A basic group box

Notice that the `GroupBox` still requires a layout container (such as a `StackPanel`) to arrange its contents. The `GroupBox` is often used to group small sets of related controls, such as radio buttons. However, the `GroupBox` has no built-in functionality, so you can use it however you want. (`RadioButton` objects are grouped by placing them into any panel. A `GroupBox` is not required, unless you want the rounded, titled border.)

## The TabItem


The `TabItem` represents a page in a `TabControl`. The only significant member that the `TabItem` class adds is the `IsSelected` property, which indicates whether the tab is currently being shown in the `TabControl`. Here's the markup that's required to create the simple example shown in Figure 6-12:

```
<TabControl Margin="5">
  <TabItem Header="Tab One">
    <StackPanel Margin="3">
      <CheckBox Margin="3">Setting One</CheckBox>
      <CheckBox Margin="3">Setting Two</CheckBox>
      <CheckBox Margin="3">Setting Three</CheckBox>
    </StackPanel>
  </TabItem>
  <TabItem Header="Tab Two">
    ...
  </TabItem>
</TabControl>
```

---

■ **Tip** You can use the `TabStripPlacement` property to make the tabs appear on the side of the tab control, rather than in their normal location at the top.

---



**Figure 6-12.** A set of tabs

As with the Content property, the Header property can accept any type of object. It displays UIElement-derived classes by rendering them and uses the ToString() method for inline text and all other objects. That means you can create a group box or a tab with graphical content or arbitrary elements in its title. Here's an example:

```
<TabControl Margin="5">
  <TabItem>
    <TabItem.Header>
      <StackPanel>
        <TextBlock Margin="3" >Image and Text Tab Title</TextBlock>
        <Image Source="happyface.jpg" Stretch="None" />
      </StackPanel>
    </TabItem.Header>

    <StackPanel Margin="3">
      <CheckBox Margin="3">Setting One</CheckBox>
      <CheckBox Margin="3">Setting Two</CheckBox>
      <CheckBox Margin="3">Setting Three</CheckBox>
    </StackPanel>
  </TabItem>

  <TabItem Header="Tab Two"></TabItem>
</TabControl>
```

Figure 6-13 shows the somewhat garish result.




**Figure 6-13.** An exotic tab title

## The Expander

The most exotic headered content control is the Expander. It wraps a region of content that the user can show or hide by clicking a small arrow button. This technique is used frequently in online help and on web pages, to allow them to include large amounts of content without overwhelming users with information they don't want to see.

Figure 6-14 shows two views of a window with three expanders. In the version on the left, all three expanders are collapsed. In the version on the right, all the regions are expanded. (Of course, users are free to expand or collapse any combination of expanders individually.)



**Figure 6-14.** Hiding content with expandable regions

Using an Expander is extremely simple—you just need to wrap the content you want to make collapsible inside. Ordinarily, each Expander begins collapsed, but you can change this in your markup (or in your code) by setting the `IsExpanded` property. Here's the markup that creates the example shown in Figure 6-14:

```
<StackPanel>
  <Expander Margin="5" Padding="5" Header="Region One">
    <Button Padding="3">Hidden Button One</Button>
  </Expander>
  <Expander Margin="5" Padding="5" Header="Region Two" >
    <TextBlock TextWrapping="Wrap">
      Lorem ipsum dolor sit amet, consectetur adipiscing elit ...
    </TextBlock>
  </Expander>
```

```

<Expander Margin="5" Padding="5" Header="Region Three">
  <Button Padding="3">Hidden Button Two</Button>
</Expander>
</StackPanel>

```

You can also choose in which direction the expander expands. In Figure 6-14, the standard value (Down) is used, but you can also set the `ExpandDirection` property to Up, Left, or Right. When the Expander is collapsed, the arrow always points in the direction where it will expand.

Life gets a little interesting when using different `ExpandDirection` values, because the effect on the rest of your user interface depends on the type of container. Some containers, such as the `WrapPanel`, simply bump other elements out of the way. Others, such as `Grid`, have the option of using proportional or automatic sizing. Figure 6-15 shows an example with a four-cell grid in various degrees of expansion. In each cell is an Expander with a different `ExpandDirection`. The columns are sized proportionately, which forces the text in the Expander to wrap. (An autosized column would simply stretch to fit the text, making it larger than the window.) The rows are set to automatic sizing, so they expand to fit the extra content.



**Figure 6-15.** Expanding in different directions

The Expander is a particularly nice fit in WPF because WPF encourages you to use a flowing layout model that can easily handle content areas that grow or shrink dynamically.

If you need to synchronize other controls with an Expander, you can handle the `Expanded` and `Collapsed` events. Contrary to what the naming of these events implies, they fire just *before* the content appears or disappears. This gives you a useful way to implement a lazy load. For example, if the content in an Expander is expensive to create, you might wait until it's shown to retrieve it. Or perhaps you want to update the content just before it's shown. Either way, you can react to the `Expanded` event to perform your work.

---

**Note** If you like the functionality of the Expander but aren't impressed with the built-in appearance, don't worry. Using the template system in WPF, you can completely customize the expand and collapse arrows so they match the style of the rest of your application. You'll learn how in Chapter 17.

---



Ordinarily, when you expand an `Expander`, it grows to fit its content. This may create a problem if your window isn't large enough to fit all the content when everything is expanded. You can use several strategies to handle this problem:

- Set a minimum size for the window (using `MinWidth` and `MinHeight`) to make sure it will fit everything even at its smallest.
- Set the `SizeToContent` property of the window so that it expands automatically to fit the exact dimensions you need when you open or close an `Expander`. Ordinarily, `SizeToContent` is set to `Manual`, but you can use `Width` or `Height` to make it expand or contract in either dimension to accommodate its content.
- Limit the size of the `Expander` by hard-coding its `Height` and `Width`. Unfortunately, this is likely to truncate the content that's inside if it's too large.
- Create a scrollable expandable region using the `ScrollViewer`.

For the most part, these techniques are quite straightforward. The only one that requires any further exploration is the combination of an `Expander` and a `ScrollViewer`. In order for this approach to work, you need to hard-code the size for the `ScrollViewer`. Otherwise, it will simply expand to fit its content. Here's an example:

```
<Expander Margin="5" Padding="5" Header="Region Two">
  <ScrollViewer Height="50">
    <TextBlock TextWrapping="Wrap">
      ...
    </TextBlock>
  </ScrollViewer>
</Expander>
```

It would be nice to have a system in which an `Expander` could set the size of its content region based on the available space in a window. However, this would present obvious complexities. (For example, how would space be shared between multiple regions when an `Expander` expands?) The `Grid` layout container might seem like a potential solution, but unfortunately, it doesn't integrate well with the `Expander`. If you try it out, you'll end up with oddly spaced rows that don't update their heights properly when an `Expander` is collapsed.

## Text Controls

WPF includes three text-entry controls: `TextBox`, `RichTextBox`, and `PasswordBox`. The `PasswordBox` derives directly from `Control`. The `TextBox` and `RichTextBox` controls go through another level and derive from `TextBoxBase`.

Unlike the content controls you've seen, the text boxes are limited in the type of content they can contain. The `TextBox` always stores a string (provided by the `Text` property). The `PasswordBox` also deals with string content (provided by the `Password` property), although it uses a `SecureString` internally to mitigate against certain types of attacks. Only the `RichTextBox` has the ability to store more sophisticated content: a `FlowDocument` that can contain a complex combination of elements.

In the following sections, you'll consider the core features of the `TextBox`. You'll end by taking a quick look at the security features of the `PasswordBox`.

---

■ **Note** The `RichTextBox` is an advanced control design for displaying `FlowDocument` objects. You'll learn how to use it when you tackle documents in Chapter 28.

---

## Multiple Lines of Text

Ordinarily, the `TextBox` control stores a single line of text. (You can limit the allowed number of characters by setting the `MaxLength` property.) However, there are many cases when you'll want to create a multiline text box for dealing with large amounts of content. In this case, set the `TextWrapping` property to `Wrap` or `WrapWithOverflow`. `Wrap` always breaks at the edge of the control, even if it means severing an extremely long word in two. `WrapWithOverflow` allows some lines to stretch beyond the right edge if the line-break algorithm can't find a suitable place (such as a space or a hyphen) to break the line.

To actually see multiple lines in a text box, it needs to be sized large enough. Rather than setting a hard-coded height (which won't adapt to different font sizes and may cause layout problems), you can use the handy `MinLines` and `MaxLines` properties. `MinLines` is the minimum number of lines that must be visible in the text box. For example, if `MinLines` is 2, the text box will grow to be at least two lines tall. If its container doesn't have enough room, part of the text box may be clipped. `MaxLines` sets the maximum number of lines that will be displayed. Even if a text box expands to fit its container (for example, a proportionally sized Grid row or the last element in a `DockPanel`), it won't grow beyond this limit.

---

■ **Note** The `MinLines` and `MaxLines` properties have no effect on the amount of content you can place in a text box. They simply help you size the text box. In your code, you can examine the `LineCount` property to find out exactly how many lines are in a text box.

---

If your text box supports wrapping, the odds are good that the user can enter more text that can be displayed at once in the visible lines. For this reason, it usually makes sense to add an always-visible or on-demand scroll bar by setting the `VerticalScrollBarVisibility` property to `Visible` or `Auto`. (You can also set the `HorizontalScrollBarVisibility` property to show a less common horizontal scroll bar.)

You may want to allow the user to enter hard returns in a multiline text box by pressing the `Enter` key. (Ordinarily, pressing the `Enter` key in a text box triggers the default button.) To make sure a text box supports the `Enter` key, set `AcceptsReturn` to `true`. You can also set `AcceptsTab` to allow the user to insert tabs. Otherwise, the `Tab` key moves to the next focusable control in the tab sequence.

---

■ **Tip** The `TextBox` class also includes a host of methods that let you move through the text content programmatically in small or large steps. They include `LineUp()`, `LineDown()`, `PageUp()`, `PageDown()`, `ScrollToHome()`, `ScrollToEnd()`, and `ScrollToLine()`.

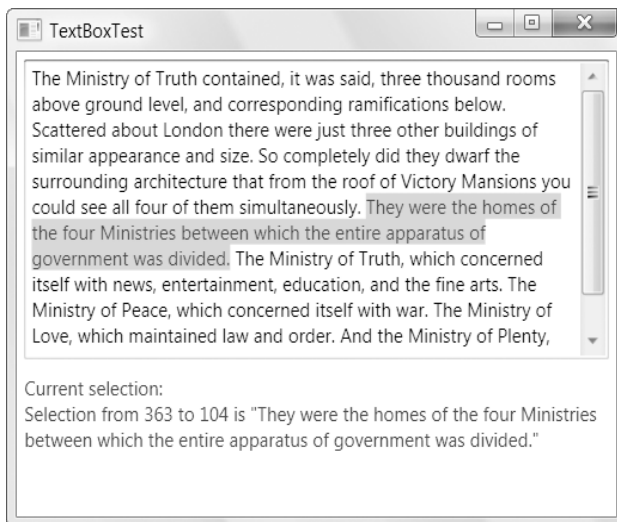
---

Sometimes, you'll create a text box purely for the purpose of displaying text. In this case, set the `IsReadOnly` property to true to prevent editing. This is preferable to disabling the text box by setting `IsEnabled` to false because a disabled text box shows grayed-out text (which is more difficult to read), does not support selection (or copying to the clipboard), and does not support scrolling.

## Text Selection

As you already know, you can select text in any text box by clicking and dragging with the mouse or holding down Shift while you move through the text with the arrow keys. The `TextBox` class also gives you the ability to determine or change the currently selected text programmatically, using the `SelectionStart`, `SelectionLength`, and `SelectedText` properties.

`SelectionStart` identifies the zero-based position where the selection begins. For example, if you set this property to 10, the first selected character is the eleventh character in the text box. The `SelectionLength` indicates the total number of selected characters. (A value of 0 indicates no selected characters.) Finally, the `SelectedText` property allows you to quickly examine or change the selected text in the text box. You can react to the selection being changed by handling the `SelectionChanged` event. Figure 6-16 shows an example that reacts to this event and displays the current selection information.



**Figure 6-16.** *Selecting text*

The `TextBox` class also includes one property that lets you control its selection behavior: `AutoWordSelection`. If this is true, the text box selects entire words at a time as you drag through the text.

Another useful feature of the `TextBox` control is `Undo`, which allows the user to reverse recent changes. The `Undo` feature is available programmatically (using the `Undo()` method), and it's available using the `Ctrl+Z` keyboard shortcut, as long as the `CanUndo` property has not been set to false.

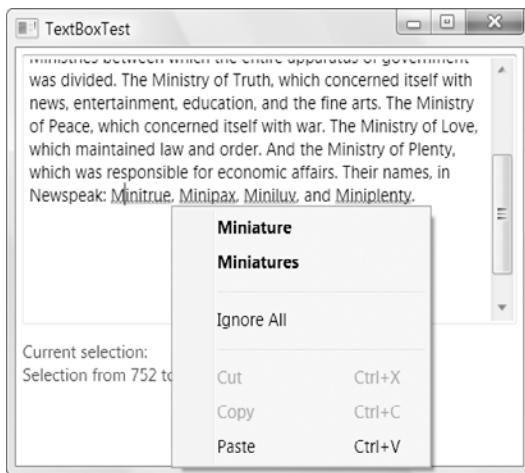
---

■ **Tip** When manipulating text in the text box programmatically, you can use the `BeginChange()` and `EndChange()` methods to bracket a series of actions that the `TextBox` will treat as a single block of changes. These actions can then be undone in a single step.

---

## Spell Checking

The `TextBox` includes an unusual frill: an integrated spell-check feature, which underlines unrecognized words with a red squiggly line. The user can right-click an unrecognized word and choose from a list of possibilities, as shown in Figure 6-17.



**Figure 6-17.** Spell-checking a text box

To turn on the spell-check functionality for the `TextBox` control, you simply need to set the `SpellCheck.IsEnabled` dependency property, as shown here:

```
<TextBox SpellCheck.IsEnabled="True">...</TextBox>
```

The spelling checker is WPF-specific and doesn't depend on any other software (such as Office). The spelling checker determines which dictionary to use based on the input language that's configured for the keyboard. You can override this default by setting the `Language` property of the `TextBox`, which is inherited from the `FrameworkElement` class, or you can set the `xml:lang` attribute on the `<TextBox>` element. However, the WPF spelling checker is currently limited to just four languages: English, Spanish, French, and German. You can use the `SpellingReform` property to set whether post-1990 spelling rule changes are applied to French and German languages.

In previous versions of WPF, the spelling checker did not support customization. WPF 4 allows you to add a list of words that will not be treated as errors (and will be used as right-click suggestions, when appropriate). To do so, you must first create a lexicon file, which is nothing more than a text file with the

extension .lex. In the lexicon file, you add the list of words. Place each word on a separate line, in any order, as shown here:

```
acantholysis
atypia
bulla
chromonychia
dermatoscopy
desquamation
...
```

In this example, the words are used regardless of the current language setting. However, you can specify that a lexicon should be used only for a specific language by adding a locale ID. Here's how you would specify that the custom words should be used only when the current language is English:

```
#LID 1033
acantholysis
atypia
bulla
chromonychia
dermatoscopy
desquamation
...
```

The other supported locale IDs are 3082 (Spanish), 1036 (French), and 1031 (German).

---

**Note** The custom dictionary feature is not designed to allow you to use additional languages. Instead, it simply augments an already supported language (like English) with the words you supply. For example, you can use a custom dictionary to recognize proper names or to allow medical terms in a medical application.

---

Once you've created the lexicon file, make sure the `SpellCheck.IsEnabled` property is set to true for your `TextBox`. The final step is to attach a `Uri` object that points to your custom dictionary, using the `SpellCheck.CustomDictionaries` property. If you choose to specify it in XAML, as in the following example, you must first import the `System` namespace so that you can declare a `Uri` object in markup:

```
<Window xmlns:sys="clr-namespace:System;assembly=system" ... >
```

You can use multiple custom dictionaries at once, as long as you add a `Uri` object for each one. Each `Uri` can use a hard-coded path to the file on a local drive or network share. But the safest approach is to use an application resource. For example, if you've added the file `CustomWords.lex` to a project named `SpellTest`, and you've set the Build Action of that file to `Resource` (using the Solution Explorer), you will use markup like this:

```
<TextBox TextWrapping="Wrap" SpellCheck.IsEnabled="True"
  Text="Now the spell checker recognizes acantholysis and offers the right correction
  for acantholysi">
  <SpellCheck.CustomDictionaries>
    <sys:Uri>pack://application:,,,/SpellTest;component/CustomWords.lex</sys:Uri>
  </SpellCheck.CustomDictionaries>
</TextBox>
```

The odd `pack://application:,,,/` portion at the beginning of the URI is the pack URI syntax that WPF uses to refer to an assembly resource. You'll take a closer look at it when you consider resources in detail in Chapter 7.

If you need to load the lexicon file from the application directory, the easiest option is to create the URI you need using code, and add it to the `SpellCheck.CustomDictionaries` collection when the window is initialized.

## The PasswordBox

The `PasswordBox` looks like a `TextBox`, but it displays a string of circle symbols to mask the characters it shows. (You can choose a different mask character by setting the `PasswordChar` property.) Additionally, the `PasswordBox` does not support the clipboard, so you can't copy the text inside.

Compared to the `TextBox` class, the `PasswordBox` has a much simpler, stripped-down interface. Much like the `TextBox` class, it provides a `MaxLength` property; `Clear()`, `Paste()` and `SelectAll()` methods; and an event that fires when the text is changed (named `PasswordChanged`). But that's it. Still, the most important difference between the `TextBox` and the `PasswordBox` is on the inside. Although you can set text and read it as an ordinary string using the `Password` property, internally the `PasswordBox` uses a `System.Security.SecureString` object exclusively.

A `SecureString` is a text-only object much like the ordinary string. The difference is how it's stored in memory. A `SecureString` is stored in memory in an encrypted form. The key that's used to encrypt the string is generated randomly and stored in a portion of memory that's never written to disk. The end result is that even if your computer crashes, malicious users won't be able to examine the paging file to retrieve the password data. At best, they will find the encrypted form.

The `SecureString` class also includes on-demand disposal. When you call `SecureString.Dispose()`, the in-memory password data is overwritten. This guarantees that all password information has been wiped out of memory and is no longer subject to any kind of exploit. As you would expect, the `PasswordBox` is conscientious enough to call `Dispose()` on the `SecureString` that it stores internally when the control is destroyed.

## List Controls

WPF includes many controls that wrap a collection of items, ranging from the simple `ListBox` and `ComboBox` that you'll examine here to more specialized controls such as the `ListView`, the `TreeView`, and the `ToolBar`, which are covered in future chapters. All of these controls derive from the `ItemsControl` class (which itself derives from `Control`).

The `ItemsControl` class fills in the basic plumbing that's used by all list-based controls. Notably, it gives you two ways to fill the list of items. The most straightforward approach is to add them directly to the `Items` collection, using code or XAML. However, in WPF, it's more common to use data binding. In this case, you set the `ItemsSource` property to the object that has the collection of data items you want to display. (You'll learn more about data binding with a list in Chapter 19.)

The class hierarchy that leads from `ItemsControls` is a bit tangled. One major branch is the *selectors*, which includes the `ListBox`, the `ComboBox`, and the `TabControl`. These controls derive from `Selector` and have properties that let you track down the currently selected item (`SelectedItem`) or its position (`SelectedIndex`). Separate from these are controls that wrap lists of items but don't support selection in the same way. These include the classes for menus, toolbars, and trees—all of which are `ItemsControls` but aren't selectors.

In order to unlock most of the features of any `ItemsControl`, you'll need to use data binding. This is true even if you aren't fetching your data from a database or an external data source. WPF data binding is general enough to work with data in a variety of forms, including custom data objects and collections. But you won't consider the details of data binding just yet. For now, you'll take only a quick look at the `ListBox` and `ComboBox`.

## The ListBox

The `ListBox` class represents a common staple of Windows design—the variable-length list that allows the user to select an item.

---

**Note** The `ListBox` class also allows multiple selection if you set the `SelectionMode` property to `Multiple` or `Extended`. In `Multiple` mode, you can select or deselect any item by clicking it. In `Extended` mode, you need to hold down the `Ctrl` key to select additional items or the `Shift` key to select a range of items. In either type of multiple-selection list, you use the `SelectedItems` collection instead of the `SelectedItem` property to get all the selected items.

---

To add items to the `ListBox`, you can nest `ListBoxItem` elements inside the `ListBox` element. For example, here's a `ListBox` that contains a list of colors:

```
<ListBox>
  <ListBoxItem>Green</ListBoxItem>
  <ListBoxItem>Blue</ListBoxItem>
  <ListBoxItem>Yellow</ListBoxItem>
  <ListBoxItem>Red</ListBoxItem>
</ListBox>
```

As you'll remember from Chapter 2, different controls treat their nested content in different ways. The `ListBox` stores each nested object in its `Items` collection.

The `ListBox` is a remarkably flexible control. Not only can it hold `ListBoxItem` objects, but it can also host any arbitrary element. This works because the `ListBoxItem` class derives from `ContentControl`, which gives it the ability to hold a single piece of nested content. If that piece of content is a `UIElement`-derived class, it will be rendered in the `ListBox`. If it's some other type of object, the `ListBoxItem` will call `ToString()` and display the resulting text.

For example, if you decided you want to create a list with images, you could create markup like this:

```
<ListBox>
  <ListBoxItem>
    <Image Source="happyface.jpg"></Image>
  </ListBoxItem>
  <ListBoxItem>
    <Image Source="happyface.jpg"></Image>
  </ListBoxItem>
</ListBox>
```

The `ListBox` is actually intelligent enough to create the `ListBoxItem` objects it needs implicitly. That means you can place your objects directly inside the `ListBox` element. Here's a more ambitious example that uses nested `StackPanel` objects to combine text and image content:

```
<ListBox>
  <StackPanel Orientation="Horizontal">
    <Image Source="happyface.jpg" Width="30" Height="30"></Image>
    <Label VerticalContentAlignment="Center">A happy face</Label>
  </StackPanel>
  <StackPanel Orientation="Horizontal">
    <Image Source="redx.jpg" Width="30" Height="30"></Image>
    <Label VerticalContentAlignment="Center">A warning sign</Label>
  </StackPanel>
  <StackPanel Orientation="Horizontal">
    <Image Source="happyface.jpg" Width="30" Height="30"></Image>
    <Label VerticalContentAlignment="Center">A happy face</Label>
  </StackPanel>
</ListBox>
```

In this example, the `StackPanel` becomes the item that's wrapped by the `ListBoxItem`. This markup creates the rich list shown in Figure 6-18.



**Figure 6-18.** *A list of images*

---

**Note** One flaw in the current design is that the text color doesn't change when the item is selected. This isn't ideal because it's difficult to read the black text with a blue background. To solve this problem, you need to use a data template, as described in Chapter 20.

---

This ability to nest arbitrary elements inside list box items allows you to create a variety of list-based controls without needing to use other classes. For example, the Windows Forms toolkit includes a



CheckedListBox class that's displayed as a list with a check box next to every item. No such specialized class is required in WPF because you can quickly build one using the standard ListBox:

```
<ListBox Name="lst" SelectionChanged="lst_SelectionChanged"
  CheckBox.Click="lst_SelectionChanged">
  <CheckBox Margin="3">Option 1</CheckBox>
  <CheckBox Margin="3">Option 2</CheckBox>
</ListBox>
```

There's one caveat to be aware of when you use a list with different elements inside. When you read the SelectedItem value (and the SelectedItems and Items collections), you won't see ListBoxItem objects; instead, you'll see whatever objects you placed in the list. In the CheckedListBox example, that means SelectedItem provides a CheckBox object.

For example, here's some code that reacts when the SelectionChanged event fires. It then gets the currently selected CheckBox and displays whether that item has been checked:

```
private void lst_SelectionChanged(object sender, SelectionChangedEventArgs e)
{
    if (lst.SelectedItem == null) return;
    txtSelection.Text = String.Format(
        "You chose item at position {0}.\r\nChecked state is {1}.",
        lst.SelectedIndex,
        ((CheckBox)lst.SelectedItem).IsChecked);
}
```

---

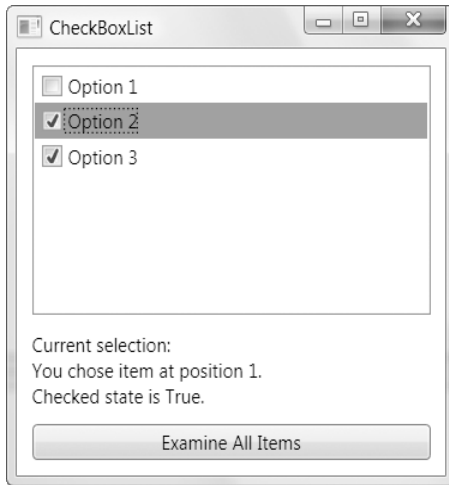
■ **Tip** If you want to find the current selection, you can read it directly from the SelectedItem or SelectedItems property, as shown here. If you want to determine which item (if any) was *unselected*, you can use the RemovedItems property of the SelectionChangedEventArgs object. Similarly, the AddedItems property tells you which items were added to the selection. In single-selection mode, one item is always added and one item is always removed whenever the selection changes. In multiple or extended mode, this isn't necessarily the case.

---

In the following code snippet, similar code loops through the collection of items to determine which ones are checked. (You could write similar code that loops through the collection of selected items in a multiple-selection list with check boxes.)

```
private void cmd_ExamineAllItems(object sender, RoutedEventArgs e)
{
    StringBuilder sb = new StringBuilder();
    foreach (CheckBox item in lst.Items)
    {
        if (item.IsChecked == true)
        {
            sb.Append(item.Content);
            sb.Append(" is checked.");
            sb.Append("\r\n");
        }
    }
    txtSelection.Text = sb.ToString();
}
```

Figure 6-19 shows the list box that uses this code.



**Figure 6-19.** A check box list

When manually placing items in a list, it's up to you whether you want to insert the items directly or explicitly wrap each one in a `ListBoxItem` object. The second approach is often cleaner, albeit more tedious. The most important consideration is to be consistent. For example, if you place `StackPanel` objects in your list, the `ListBox.SelectedItem` object will be a `StackPanel`. If you place `StackPanel` objects wrapped by `ListBoxItem` objects, the `ListBox.SelectedItem` object will be a `ListBoxItem`, so code accordingly.

The `ListBoxItem` offers a little extra functionality from what you get with directly nested objects. Namely, it defines an `IsSelected` property that you can read (or set) and a `Selected` and `Unselected` event that tells you when that item is highlighted. However, you can get similar functionality using the members of the `ListBox` class, such as the `SelectedItem` (or `SelectedItems`) property, and the `SelectionChanged` event.

Interestingly, there's a technique to retrieve a `ListBoxItem` wrapper for a specific object when you use the nested object approach. The trick is the often overlooked `ContainerFromElement()` method. Here's the code that checks whether the first item is selected in a list using this technique:

```
ListBoxItem item = (ListBoxItem)lst.ContainerFromElement(
    (DependencyObject)lst.SelectedItems[0]);
MessageBox.Show("IsSelected: " + item.IsSelected.ToString());
```

## The ComboBox

The `ComboBox` is similar to the `ListBox` control. It holds a collection of `ComboBoxItem` objects, which are created either implicitly or explicitly. As with the `ListBoxItem`, the `ComboBoxItem` is a content control that can contain any nested element.

The key difference between the `ComboBox` and `ListBox` classes is the way they render themselves in a window. The `ComboBox` control uses a drop-down list, which means only one item can be selected at a time.

If you want to allow the user to type text in the combo box to select an item, you must set the `IsEditable` property to true, and you must make sure you are storing ordinary text-only `ComboBoxItem` objects or an object that provides a meaningful `ToString()` representation. For example, if you fill an editable combo box with `Image` objects, the text that appears in the upper portion is simply the fully qualified `Image` class name, which isn't much use.

One limitation of the `ComboBox` is the way it sizes itself when you use automatic sizing. The `ComboBox` widens itself to fit its content, which means that it changes size as you move from one item to the next. Unfortunately, there's no easy way to tell the `ComboBox` to take the size of its largest contained item. Instead, you may need to supply a hard-coded value for the `Width` property, which isn't ideal.

## Range-Based Controls

WPF includes three controls that use the concept of a *range*. These controls take a numeric value that falls in between a specific minimum and maximum value. These controls—`ScrollBar`, `ProgressBar`, and `Slider`—all derive from the `RangeBase` class (which itself derives from the `Control` class). But although they share an abstraction (the range), they work quite differently.

The `RangeBase` class defines the properties shown in Table 6-4.

**Table 6-4.** *Properties of the RangeBase Class*

Name	Description
Value	The current value of the control (which must fall between the minimum and maximum). By default, it starts at 0. Contrary to what you might expect, <code>Value</code> isn't an integer—it's a double, so it accepts fractional values. You can react to the <code>ValueChanged</code> event if you want to be notified when the value is changed.
Maximum	The upper limit (the largest allowed value).
Minimum	The lower limit (the smallest allowed value).
SmallChange	The amount the <code>Value</code> property is adjusted up or down for a small change. The meaning of a "small change" depends on the control (and may not be used at all). For the <code>ScrollBar</code> and <code>Slider</code> , this is the amount the value changes when you use the arrow keys. For the <code>ScrollBar</code> , you can also use the arrow buttons at either end of the bar.
LargeChange	The amount the <code>Value</code> property is adjusted up or down for a large change. The meaning of a "large change" depends on the control (and may not be used at all). For the <code>ScrollBar</code> and <code>Slider</code> , this is the amount the value changes when you use the Page Up and Page Down keys or when you click the bar on either side of the thumb (which indicates the current position).

Ordinarily, there's no need to use the ScrollBar control directly. The higher-level ScrollViewer control, which wraps two ScrollBar controls, is typically much more useful. The Slider and ProgressBar are more practical, and are often useful on their own.

## The Slider

The Slider is a specialized control that's occasionally useful—for example, you might use it to set numeric values in situations where the number itself isn't particularly significant. For example, it makes sense to set the volume in a media player by dragging the thumb in a slider bar from side to side. The general position of the thumb indicates the relative loudness (normal, quiet, or loud), but the underlying number has no meaning to the user.

The key Slider properties are defined in the RangeBase class. Along with these, you can use all the properties listed in Table 6-5.

**Table 6-5.** *Additional Properties in the Slider Class*

Name	Description
Orientation	Switches between a vertical and a horizontal slider.
Delay and Interval	Control how fast the thumb moves along the track when you click and hold down either side of the slider. Both are millisecond values. The Delay is the time before the thumb moves one (small change) unit after you click, and the Interval is the time before it moves again if you continue holding down the mouse button.
TickPlacement	Determines where the tick marks appear. (Tick marks are notches that appear near the bar to help you visualize the scale.) By default, TickPlacement is set to None, and no tick marks appear. If you have a horizontal slider, you can place the tick marks above (TopLeft) or below (BottomRight) the track. With a vertical slider, you can place them on the left (TopLeft) and right (BottomRight). (The TickPlacement names are a bit confusing because two values cover four possibilities, depending on the orientation of the slider.)
TickFrequency	Sets the interval in between ticks, which determines how many ticks appear. For example, you could place them every 5 numeric units, every 10, and so on.
Ticks	If you want to place ticks in specific, irregular positions, you can use the Ticks collection. Simply add one number (as a double) to this collection for each tick mark. For example, you could place ticks at the positions 1, 1.5, 2, and 10 on the scale by adding these numbers.